

# Building Efficient Recommender Systems: Techniques, Challenges, and Best Practices

**KATERINA ILIAKOPOULOU-ZANOS**

Machine Learning Software Engineer  
Meta



# Recommender Systems are everywhere

FACEBOOK

THE NEW YORK TIMES

MICROSOFT

INSTAGRAM

PINTEREST

BUZZFEED

GOOGLE

TIKTOK

AMAZON

TINDER

SPOTIFY

NETFLIX



# 15%

**of businesses' ML projects  
scale across multiple parts of  
the business.**

Source: [Operationalizing machine learning in processes](#)

# 87%


**of machine learning models  
never make it into production.**

Source: [Why do 87% of data science projects never make it into production?](#)



# Building Recommender Systems Efficiently

# What is a recommender system



# Choosing the right model

**Trending / Most Popular**

**Collaborative Filtering (User-User, Item-Item, User-Item)**

**Content Based**

**Contextual Bandits**

**Deep Learning**



# Get the data

User History / Engagement data

Precomputed features

Embeddings



# Evaluation

## Offline Evaluation

- Recall
- Precision
- F1 Score
- Calibration

## A/B Testing

- Stat-sig results on core business value metrics





# Building efficient recommender systems

1. **Responsiveness**
  2. **Speed of development**
  3. **Cost of deployment**
-

# Responsiveness



**[Responsiveness]** How fast can we fetch fresh content to recommend to the user?

Indexing content and extracting features as soon as it becomes available

Identify potential bias in the model that could punish new content

Leverage exploration / exploitation techniques to allow for gathering the right data for recommending content to the right audience

Choosing between offline & online inference




**[Responsiveness]** How fast can we deliver recommendations to the user?

Cache

Replication for handling multiple requests


Optimizing data requests





How our system needs to  
respond defines key  
architecture & modeling  
decisions

# Speed of development



**[Speed of Development]** How do you align work and priorities with your team and partner teams?

How big is your team? Are you working with ML Engineers, Data Scientists, Data Engineers?

Getting features / data ready for training the model and later for inference

Working with partner teams on conflicting goals



# **[Speed of Development]** Working with devops on architecting the desired solution

**How many systems do we have to touch in order to make our solution production ready?**

**Are the changes simple and straightforward or do we have to build new functionalities to make our solution available?**



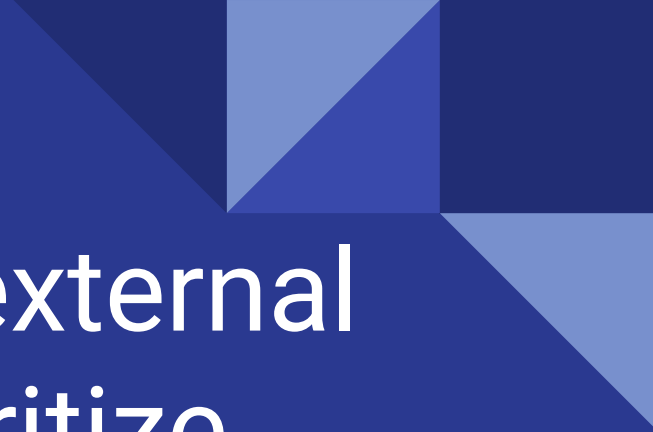


**[Speed of Development]** How can we test early and have results to know if a recommender solution is effective for our product?

**Do upfront analysis to prove that the recommender solution could work well with your data before moving to online testing**

**Test gradually. Develop an MVP with minimum features before jumping into a more complex system**





Identify internal and external dependencies to prioritize efficiently.

Test often - both offline and online.

# Cost of deployment



**[Cost of Deployment]** What are we willing to sacrifice over achieving desired success goals?

**CPU vs Storage Space**

**DB responsiveness and availability**

**Data & Model cost**

**Work hours**




**[Cost]** What is the accepted cost of my recommender system?

What do we gain with the product we are shipping?

If we prompt for a cheaper solution how much does it hurt our user value?

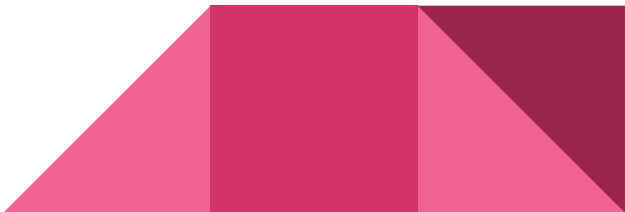
If we achieve the desired trade off now, how fast is this bound to change as our user base expands?





If possible, decide on what your cost ceiling is first. It will set the parameters of your system later.

# Building efficient recommender systems

1. **Decide how the recommender system drives user value.**
  2. **Set an upper limit for system cost.**
  3. **Make architecture decisions for the system and model based on how responsive we want it to be in terms of content delivery and freshness.**
  4. **Optimize for speed of development by identifying key dependencies and aligning on priorities with partner teams.**
  5. **Assess the cost on resources and choose the best tradeoff.**
  6. **Agree on launch criteria within your org.**
  7. **Test gradually and often as you develop the system.**
  8. **Plan for scaling the solution to a larger user base.**
- 

# Thank you!



**LinkedIn:**

<https://www.linkedin.com/in/katerinailiakopoulou/>

**Twitter: @matzika00**

<https://www.katerinazanos.com/>

